

Tweet Clustering

*Exploratory data analysis of one million tweets using clustering techniques in scikit-learn.
Civil & Environmental Engineering 263n: Scalable Spatial Analytics at UC-Berkeley, Fall 2016
By Paul Sohn, September 14, 2016*

The scikit-learn package in Python includes many options for clustering algorithms. We will get to know three of them using a dataset of 1 million tweets (mostly) from the Bay Area. We will start with a random subset of 100,000 tweets to test the algorithms. The clustering algorithms we are using here are:

- K-means
- MiniBatch K-means
- DBSCAN

The first step is to understand the processing limits of each algorithm.

Part 1: Baseline Results for Different Clustering Algorithms

K-means

We are trying to find the reference time of clustering of 100K samples into $k=100$ clusters with k-means. The basic python code snippet involves instantiating a KMeans object, fitting to a numpy array (`data`), and printing the time taken:

```
k_means = KMeans(n_clusters=100, init='k-means++', n_init=10)

t0 = time.time()
k_means.fit(data)
print time.time() - t0
```

Time to cluster 100,000 tweets into 100 clusters using K-means: **20.7 seconds**

We can implement the above processing time test using various values for the number of requested clusters k , increasing k until processing time reaches an arbitrary threshold above which performance is no longer acceptable (code for this loop is below in Part 2). If we select 60 seconds as the threshold, we find a **k_{\max}** of **~311 clusters**.

MiniBatch K-means

As above, we are trying to find the reference time of clustering of 100K samples into $k=100$ clusters with MiniBatch k-means. The python code is very similar, except one has to select a batch size. We can simply test several arbitrary batch size values as below:

```
for batch_size in [5, 10, 20, 50, 100, 500, 1000]:
```

```
    mb = MiniBatchKMeans(n_clusters=100, init='k-means++', n_init=10, batch_size=batch_size)

    t0 = time.time()
    mb.fit(data)
    print time.time() - t0
```

Batch Size	Time to generate 100 clusters (seconds)
5	5.29
10	2.84

Batch Size	Time to generate 100 clusters (seconds)
20	2.69
50	1.48
100	0.74
500	0.63
1000	0.70

Like with k-means above, we can find `k_max` for various batch sizes (code for this loop is below in Part 2). If we select 60 seconds as the threshold, we find the following values of `k_max`:

Batch Size	Maximum Number of Clusters
5	5000
10	5300
20	4950
50	5200

Interestingly, it turns out that the maximum number of clusters does not vary with batch size, but in any case, `k_max` is much, much higher than for the basic K-means algorithm.

DBSCAN

DBSCAN is different from the other algorithms in that it does not produce a set number of clusters but instead detects as many clusters as exist based on two parameters:

- `eps`, or the “The maximum distance between two samples for them to be considered as in the same neighborhood” and
- `min_samples`, the “number of samples (or total weight) in a neighborhood for a point to be considered as a core point.” (Language taken from sklearn documentation).

Our goal here is to find the value of `eps` in DBScan resulting in approximately 100 clusters (`eps_100`) of a minimum of samples (`min_samples=100`) and the corresponding processing time. Furthermore, we can convert `eps`, which is in degrees latitude and longitude, into a more meaningful unit like miles. We will use a *very* rough approximation of 100 kilometers per degree:

```
miles = .75
kilometers = miles / 0.621371
eps = kilometers / 100
```

Then we can loop through various values of `eps` to find `eps_100`, the value that will give us approximately 100 clusters:

```
dbscan = DBSCAN(eps=eps, min_samples=100)
dbscan.fit(data)
print len(numpy.unique(dbscan.labels_))
```

miles	epsilon	clusters	seconds
0.1	0.00160934449789	45	0.873108863831
0.2	0.00321868899579	70	1.11532998085
0.3	0.00482803349368	106	1.34408187866
0.4	0.00643737799157	126	1.49612116814
0.5	0.00804672248946	119	1.59469985962
0.6	0.00965606698736	103	1.79232883453

miles	epsilon	clusters	seconds
0.7	0.0112654114852	86	2.09918904305
0.8	0.0128747559831	72	2.40508508682
0.9	0.014484100481	61	5.46035599709

I select **0.009656** as `eps_100`. I could have also used a smaller value but chose to be inclusive.

Part 2: Scalability of Clustering Algorithms

In this section, we will expand the baseline results from the previous section to see how performance scales by:

- Number of clusters (K-means and MiniBatch K-means only)
- Number of data points processed

K-means

Number of requested clusters

We can find `k_max` as reported above by looping the K-means model through values of `k` (number of clusters) with the set of 100,000 tweets:

```
for k in range(2, 10000, 5): # I externally set a timeout of 60 seconds
    k_means = KMeans(n_clusters=k, init='k-means++', n_init=10)

    t0 = time.time()
    k_means.fit(data)
    print time.time() - t0
```

Results are as follows (Figure 1 in PDF); processing time for the K-means algorithm seems to scale linearly at a rate of 0.181 seconds per cluster.

Number of data points processed

We are now trying to find the processing time of K-means for varying numbers of data samples (consider the range of 100 to 100,000) for a fixed `k=100`. To achieve this, we simply run a loop fitting the K-means model to random samples of various sizes:

```
for n in range(100, 100000, 500):
    k_means = KMeans(n_clusters=100, init='k-means++', n_init=10)

    data = data[numpy.random.randint(low=0, high=len(data), size=n), :]

    t0 = time.time()
    k_means.fit(data)
    print time.time() - t0
```

The results are as follows (Figure 2 in PDF). Again, there is a linear pattern of about 1.84 seconds of processing time per 10,000 additional tweets. Extrapolating this rate, generating 100 clusters from a dataset of 1,000,000 tweets would take about **184 seconds**, three times our acceptable threshold of 60 seconds.

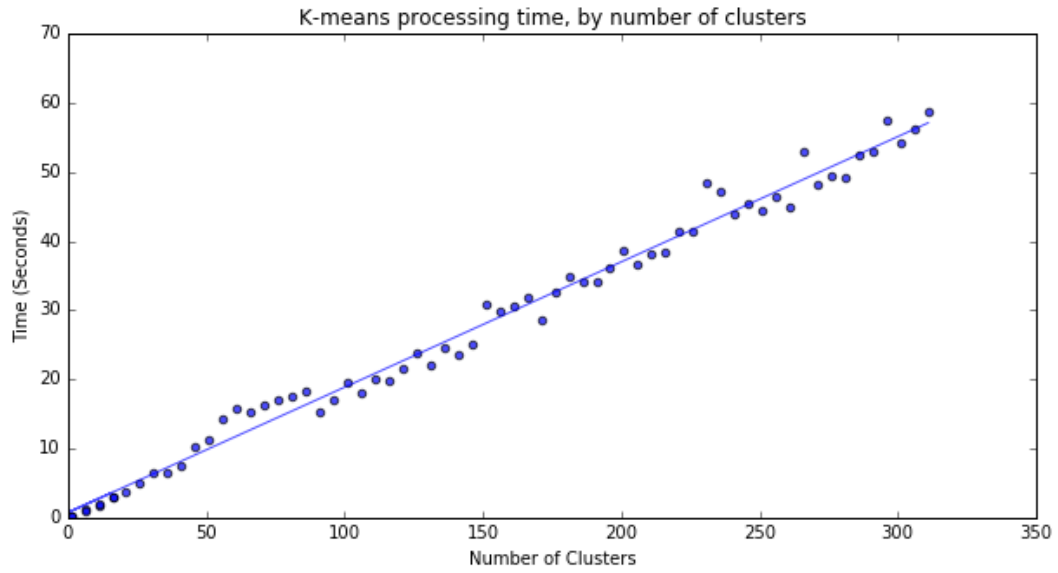


Figure 1: K-means processing time by number of clusters




Figure 2: K-means processing time by sample size

MiniBatch K-means

Number of requested clusters

We can find `k_max` similarly to above, looping through `k` (number of clusters) as well as values of `batch_size` with the set of 100,000 tweets:

```
for batch_size in [5, 10, 20, 50, 100]:
    for k in range(2, 10000, 5): # I externally set a timeout of 60 seconds
        mb = MiniBatchKMeans(n_clusters=k, init='k-means++', n_init=10, batch_size=batch_size, init_size=100000)

        t0 = time.time()
        mb.fit(data)
        print time.time() - t0
```

Results below (Figure 3) suggest that processing time scales at an increasing rate with number of clusters requested, but of course processing time is much lower than for equivalent numbers of clusters with the basic K-means algorithm.

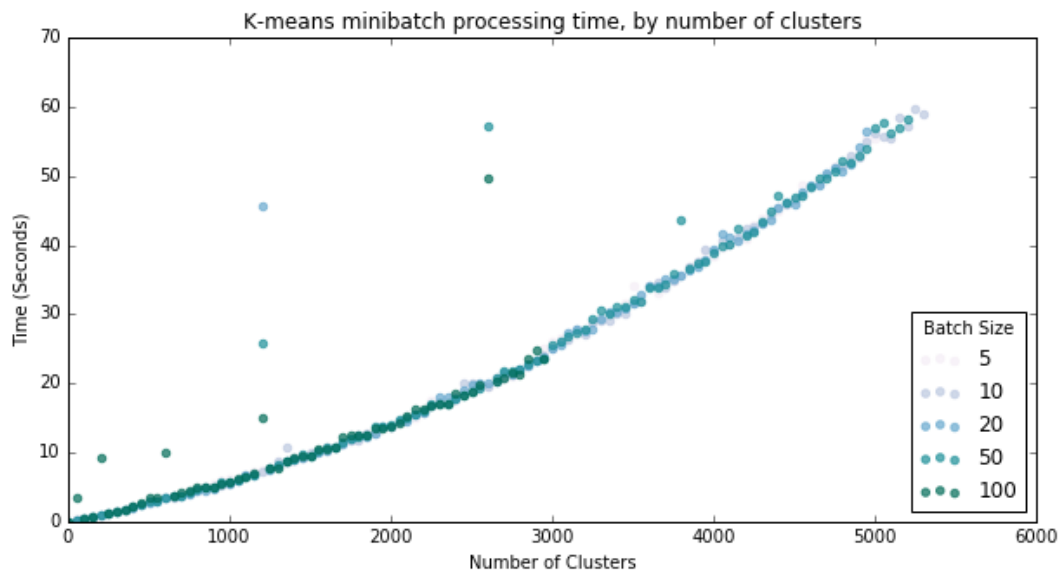


Figure 3: Minibatch processing time by number of clusters

Number of data points processed

We use a similar loop through varying sample sizes (from 100 to 100,000) at `k=100`:

```
for batch_size in [5, 10, 20, 50, 100]:
    for n in range(100, 100000, 500):
        mb = MiniBatchKMeans(n_clusters=100, init='k-means++', n_init=10, batch_size=batch_size)

        data = data[np.random.randint(low=0, high=len(data), size=n), :]

        t0 = time.time()
        mb.fit(data)
        print time.time() - t0
```

These results (Figure 4) suggest that processing time for 100 clusters essentially *does not scale* with sample size, at least in the vicinity of our sample sizes. So we can expect processing time for 1,000,000 tweets to be between **0-5 seconds**.

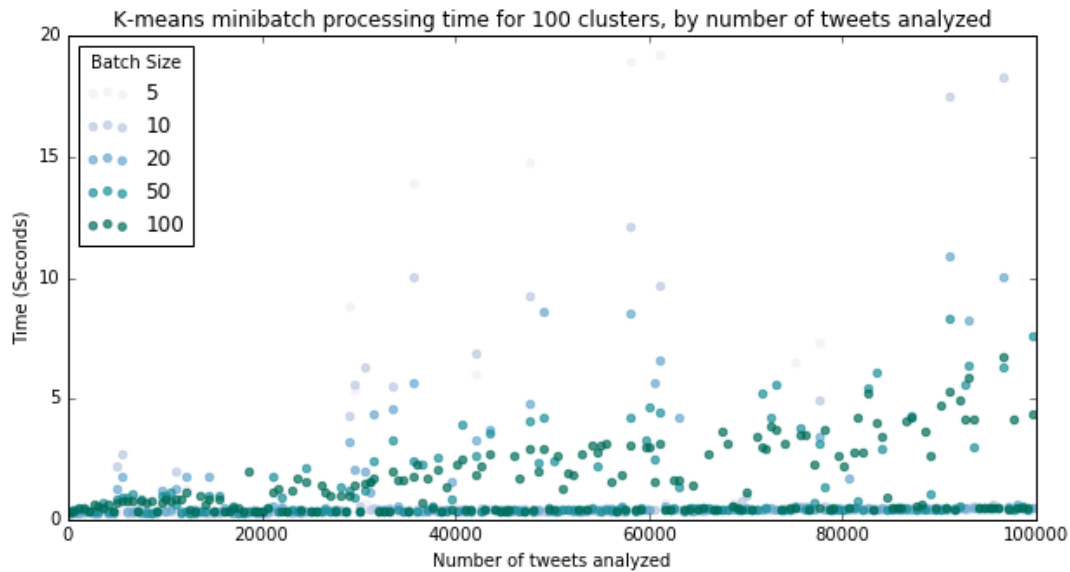


Figure 4: Minibatch K-means processing time by sample size

DBSCAN

Number of data points processed

Like the other algorithms, we will loop through various sample sizes (from $n=100$ to 100,000), this time using a fixed `eps100=0.009656` that we found above.

```
for n in range(100, 100000, 500):
    dbscan = DBSCAN(eps=eps,
                    min_samples=min_samples)

    data = data[np.random.randint(low=0, high=len(data), size=n), :]

    t0 = time.time()
    dbscan.fit(data)
    print time.time() - t0
```

I did not set a legend for the results below (Figure 5) but the size of circles is proportional to the number of clusters detected. These results show that this implementation, at this value of `eps`, does scale quite well with sample size, and should process 1,000,000 samples within **15-25 seconds**; well within our acceptable range.

Summary of scalability

We've learned a couple of things about the three algorithms and how their performance scales:

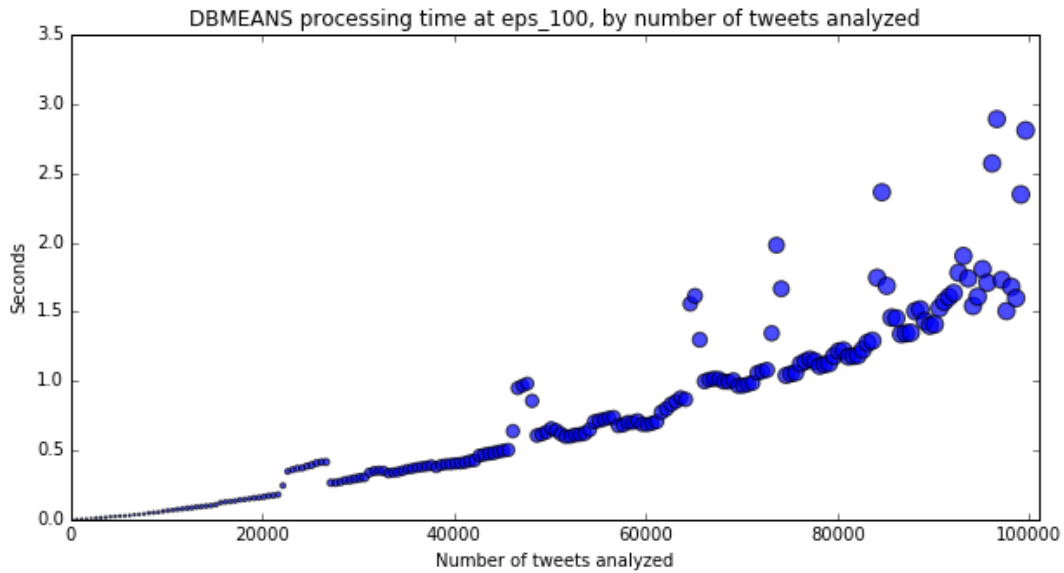


Figure 5: DBSCAN processing time by sample size

- K-means does not scale well, in terms of the size of the dataset, and in terms of detecting large numbers of clusters. Processing time scales linearly. While getting 100 clusters is not too slow for a sample of $n=100,000$, this algorithm doesn't perform well for larger datasets.
- MiniBatch K-means also does not scale well in terms of number of clusters, but is much faster than the vanilla K-means implementation, depending on batch size.
- MiniBatch K-means does scale well with larger datasets, if the number of clusters is low (in our example, $k=100$). This means we can feasibly use this to detect 100 clusters for even a dataset with 1,000,000 samples.
- DBSCAN